

# What is R?

R is a program. Like any program, it takes inputs from a user and does something.

In this case, the input is your code.

You can input code in two ways:

1. R can read it from a file
2. R can read it “interactively”

We'll focus on using R interactively in this module.

# Does Something

What do we want R to do? What is the output?

This will determine how you use R (or any other language).

# Communication

If you want to communicate something to someone:

- ▶ HTML
- ▶ PDF Document
- ▶ Jupyter Notebook
- ▶ Website
- ▶ Pretty Pictures

Plus the code to repeat/modify the above material

# ETL

If you want to transform data:

- ▶ CSV, JSON, RData, SQL, database entries

Plus the code to repeat/modify the transformation

# Your Own Opinion

If you want to learn, ask questions, the output might be:

- ▶ Nothing

Plus the code to repeat/modify your question or answer.

# You're writing code

No matter what you use R for, the code is always an end in and of itself.

# R is functional

You can manage code complexity in R with functions.

1. Write functions that have one responsibility.
2. Assign each responsibility fully to one function.

# Learning R

Today we want to use R to do the following:

Test the accuracy of the standard errors in an OLS regression, given input data of different sizes, with Gaussian noise.



# Break down the problem

The first step is to break down the task into separate responsibilities:

1. Generate input data with Gaussian noise, of differing sizes.
2. Get standard errors of OLS regression on the above input data.
3. Compare OLS estimate to the true value.

# Responsibility

Given our breakdown of the tasks in this problem, we know we need 3 functions.

Let's start with the first one, generating data.

# Functions in R

Here is a function that generates zeros:

```
generate_data <- function (howMuchData) {  
  output <- rep(0, how_much_data)  
  return(output)  
}
```

How does variable assignment work? Function arguments? Function calls?

# Functions in R

Remember the return statement is optional. This is the same:

```
generate_data <- function (how_much_data) {  
  output <- rep(0, how_much_data)  
  output  
}
```

# Functions in R

What if we do this? Is this the same function?

```
how_much_data <- 20
```

```
generate_data <- function () {  
  output <- rep(0, how_much_data)  
  output  
}
```

# Functions in R

If R doesn't find a variable inside the function, it will look outside.

# Functions in R

What's bad about this?

```
generate_data <- function () {  
  output <- rep(0, how_much_data)  
  output  
}
```

```
how_much_data <- 20  
data <- generate_data()
```

```
how_much_data <- 35  
data_2 <- generate_data()
```

# Functions in R

Remember the responsibility metaphor. Previously, the function `generate_data` was completely in charge of generating data.

In this version, the responsibility of generating the data is shared: between the function and whoever is in charge of keeping track of the value of `“how_much_data”` in the global scope.



# Functions in R

That someone is you.

But you shouldn't be sharing the responsibility. You have enough to do.

Give the responsibility fully to the function. Write a reliable function.

## Functions in R

```
generate_data <- function (how_much_data) {  
  output <- rep(0, how_much_data)  
  output  
}
```

```
data <- generate_data(20)
```

```
data_2 <- generate_data(35)
```

# Generating Data

Let's work on generating data according to the data generating process:

$$y = \beta x + \epsilon$$

Where

$$x \sim \text{Normal}(0, 1)$$

$$\epsilon \sim \text{Normal}(0, \sigma^2)$$

## Generating Data

```
generate_data <- function(N, beta, sd) {  
  x <- rnorm(N, 0, 1)  
  eps <- rnorm(N, 0, sd)  
  y <- beta*x + eps  
  list(x = x, y = y)  
}
```

What are lists in R?

## Generating Data

```
generate_data <- function(N, beta, sd) {  
  x <- rnorm(N, 0, 1)  
  eps <- rnorm(N, 0, sd)  
  y <- beta*x + eps  
  list(x = x, y = y)  
}
```

What is a list?

## Running the Regression

```
run_regression <- function(y, x) {  
  coef <- ???  
  se <- ???  
  list(coef=coef, se=se)  
}
```

Here again we have two responsibilities:

1. Calculating the coefficient on  $x$ .
2. Calculating the standard errors on the coefficients.

## Running the Regression

```
calc_coef <- function(y,x) cov(x,y) / var(x)

calc_se <- function(y, x, coef) {
  n <- length(y)
  eps <- y - x*coef
  e_sd <- mean(eps^2)
  se <- sqrt(e_sd / (n*var(x)))
  se
}
```

- ▶ Functions can be put on one line!
- ▶ `var()` actually returns the unbiased variance estimate  $(n-1)$ , but we will just use it to keep things simple.

## Running the Regression

```
run_regression <- function(y, x) {  
  coef <- calc_coef(y, x)  
  se <- calc_se(y, x)  
  list(coef=coef, se=se)  
}
```



# Evaluating the Model

```
eval_model <- function(coef, se, beta, conf = 1.96) {  
  up <- coef + se*conf  
  down <- coef - se*conf  
  beta > down & beta < up  
}
```

- ▶ Default function values
- ▶ Logical operators
- ▶ Will this return a scalar or a vector? (depends on coef and beta!)

# Evaluating the Model

```
eval_model <- function(coef, se, beta, conf = 1.96) {  
  up <- coef + se*conf  
  down <- coef - se*conf  
  beta > down & beta < up  
}
```

Note:

- ▶ Default function values
- ▶ Logical operators
- ▶ Will this return a scalar or a vector?

## Making our simulation

```
simulate <- function(N, beta, sd) {  
  d <- generate_data(N, beta, sd)  
  m <- run_regression(d$y, d$x)  
  eval_model(m$coef, m$se, beta)  
}
```

```
avg_simulations <- function(M, N, beta, sd) {  
  inside <- sapply(1:M, function(x) simulate(N, beta, sd))  
  sum(inside) / M  
}
```

- ▶ apply family
- ▶ Summing boolean vector

## Results

Let's see our results:

```
library(ggplot2)

check_N <- function(M, beta, sd) {
  x <- seq(4, 50, 2)
  y <- sapply(x, function(N) avg_simulations(M, N, beta,
    qplot(x, y)
  })
}
```

# Results




Figure 1: inline